

## HOUSING PRICE IN THE STATE OF WASHINGTON

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: df = pd.read_csv('home_data.csv')
print(df.shape)
print(df.dtypes)

#Rest set the price column for avoiding crash on the regression model later
df['price_new'] = df['price']/1000
```


```
(21613, 21)
id                int64
date              object
price             int64
bedrooms          int64
bathrooms         float64
sqft_living       int64
sqft_lot          int64
floors            float64
waterfront        int64
view              int64
condition         int64
grade             int64
sqft_above        int64
sqft_basement     int64
yr_built          int64
yr_renovated      int64
zipcode           int64
lat               float64
long              float64
sqft_living15     int64
sqft_lot15        int64
dtype: object
```

In [3]: `df.head()`

Out[3]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	water
0	7129300520	20141013T000000	221900	3	1.00	1180	5650	1.0	
1	6414100192	20141209T000000	538000	3	2.25	2570	7242	2.0	
2	5631500400	20150225T000000	180000	2	1.00	770	10000	1.0	
3	2487200875	20141209T000000	604000	4	3.00	1960	5000	1.0	
4	1954400510	20150218T000000	510000	3	2.00	1680	8080	1.0	

5 rows × 22 columns



```
In [4]: df_numeric = df.select_dtypes(include='number')
numeric_cols = df_numeric.columns.values
print(numeric_cols)
print(df_numeric)
```

```
['id' 'price' 'bedrooms' 'bathrooms' 'sqft_living' 'sqft_lot' 'floors'
'waterfront' 'view' 'condition' 'grade' 'sqft_above' 'sqft_basement'
'yr_built' 'yr_renovated' 'zipcode' 'lat' 'long' 'sqft_living15'
'sqft_lot15' 'price_new']
```

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors
\							
0	7129300520	221900	3	1.00	1180	5650	1.0
1	6414100192	538000	3	2.25	2570	7242	2.0
2	5631500400	180000	2	1.00	770	10000	1.0
3	2487200875	604000	4	3.00	1960	5000	1.0
4	1954400510	510000	3	2.00	1680	8080	1.0
...	...	...	...	...	...	...	...
21608	263000018	360000	3	2.50	1530	1131	3.0
21609	6600060120	400000	4	2.50	2310	5813	2.0
21610	1523300141	402101	2	0.75	1020	1350	2.0
21611	291310100	400000	3	2.50	1600	2388	2.0
21612	1523300157	325000	2	0.75	1020	1076	2.0


	waterfront	view	condition	...	sqft_above	sqft_basement	yr_built	\
0	0	0	3	...	1180	0	1955	
1	0	0	3	...	2170	400	1951	
2	0	0	3	...	770	0	1933	
3	0	0	5	...	1050	910	1965	
4	0	0	3	...	1680	0	1987	
...	...	...	...	...	...	...	...	
21608	0	0	3	...	1530	0	2009	
21609	0	0	3	...	2310	0	2014	
21610	0	0	3	...	1020	0	2009	
21611	0	0	3	...	1600	0	2004	
21612	0	0	3	...	1020	0	2008	

	yr_renovated	zipcode	lat	long	sqft_living15	sqft_lot15	\
0	0	98178	47.5112	-122.257	1340	5650	
1	1991	98125	47.7210	-122.319	1690	7639	
2	0	98028	47.7379	-122.233	2720	8062	
3	0	98136	47.5208	-122.393	1360	5000	
4	0	98074	47.6168	-122.045	1800	7503	
...	...	...	...	...	...	...	
21608	0	98103	47.6993	-122.346	1530	1509	
21609	0	98146	47.5107	-122.362	1830	7200	
21610	0	98144	47.5944	-122.299	1020	2007	
21611	0	98027	47.5345	-122.069	1410	1287	
21612	0	98144	47.5941	-122.299	1020	1357	

	price_new
0	221.900
1	538.000
2	180.000
3	604.000
4	510.000
...	...
21608	360.000

```
21609    400.000
21610    402.101
21611    400.000
21612    325.000
```

```
[21613 rows x 21 columns]
```




```
In [5]: print(df.isnull().sum())
```

```
id            0
date          0
price         0
bedrooms      0
bathrooms     0
sqft_living   0
sqft_lot      0
floors        0
waterfront    0
view          0
condition     0
grade         0
sqft_above    0
sqft_basement 0
yr_built      0
yr_renovated  0
zipcode       0
lat           0
long          0
sqft_living15 0
sqft_lot15    0
price_new     0
dtype: int64
```

Checking if there is any missing values in our dataset for all variables.

```
In [6]: cols = df.columns[:]
plt.title('Heatmap of Null Values for All Variables')
colors = ['#000099', '#ffff00']
sns.heatmap(df[cols].isnull(), cmap=sns.color_palette(colors))
plt.show()
```



We are dropping the last 4 columns because we are not considering using these variables for our analysis.

```
In [7]: cols_to_drop = ['lat', 'long', 'sqft_living15', 'sqft_lot15']
df = df.drop(cols_to_drop, axis=1)
```

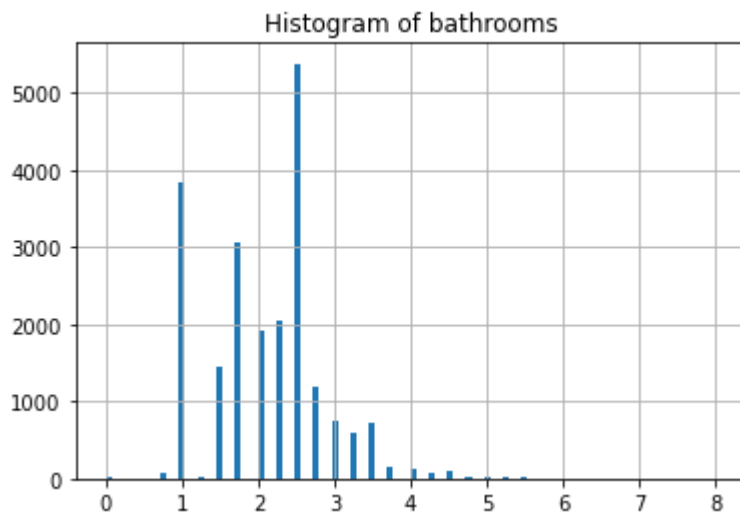
Chosen dependent variable as 'price\_new', independent variables as 'grade', 'sqft\_living', 'floors', & 'bathrooms'.

# Histogram


```
In [8]: # 1. histogram of price_new. to detect outliers  
plt.title('Histogram of price_new ')  
df.price_new.hist(bins=100)  
plt.show()
```




```
In [9]: # 2. histogram of bedrooms. to detect outliers  
plt.title('Histogram of bathrooms')  
df.bathrooms.hist(bins=100)  
plt.show()
```



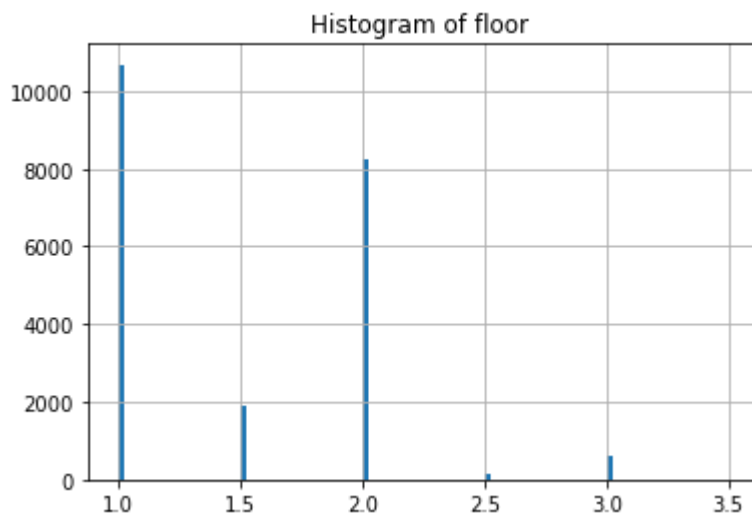
```
In [10]: # 3. histogram of sqft_living. to detect outliers
plt.title('Histogram of sqft_living ')
df.sqft_living.hist(bins=100)
plt.show()
```



```
In [11]: # 4. histogram of grade. to detect outliers
df.grade.hist(bins=100)
plt.title('Histogram of grade ')
plt.show()
```




```
In [12]: # 5. histogram of floor. to detect outliers
df.floors.hist(bins=100)
plt.title('Histogram of floor ')
plt.show()
```



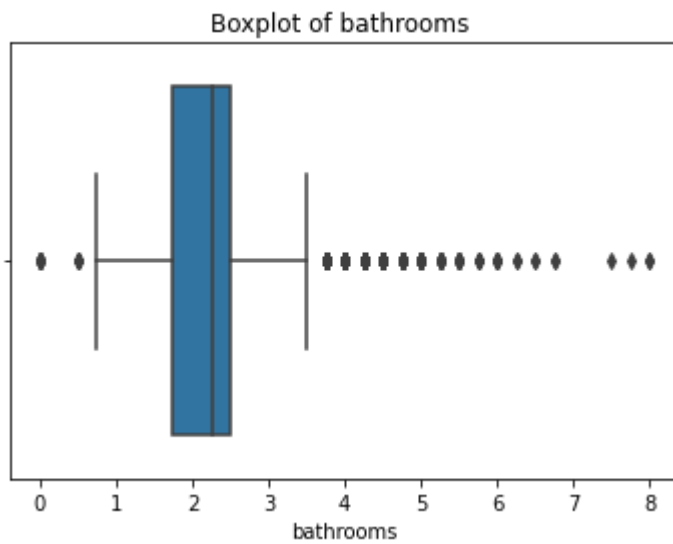
## Boxplots

```
In [13]: # 1. box plot of price_new to detect outliers
plt.title('Boxplot of price_new ')
sns.boxplot(data=df, x='price_new')
plt.show()
```

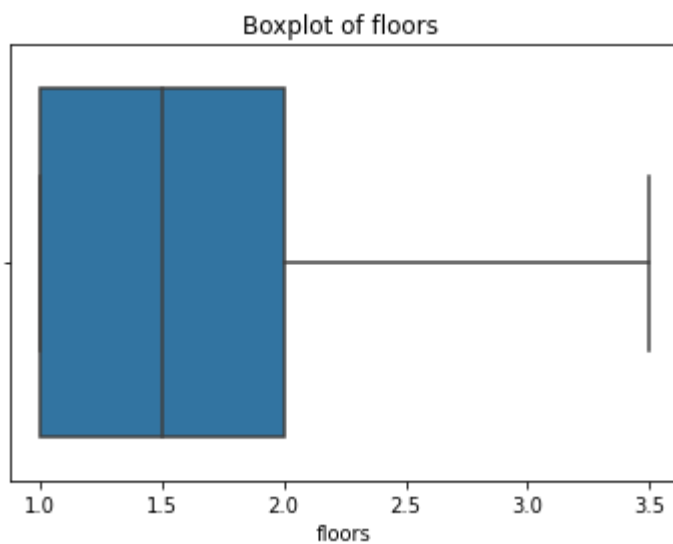





```
In [14]: # 2. box plot of bathrooms to detect outliers
plt.title('Boxplot of bathrooms ')
sns.boxplot(data=df,x='bathrooms')
plt.show()
```



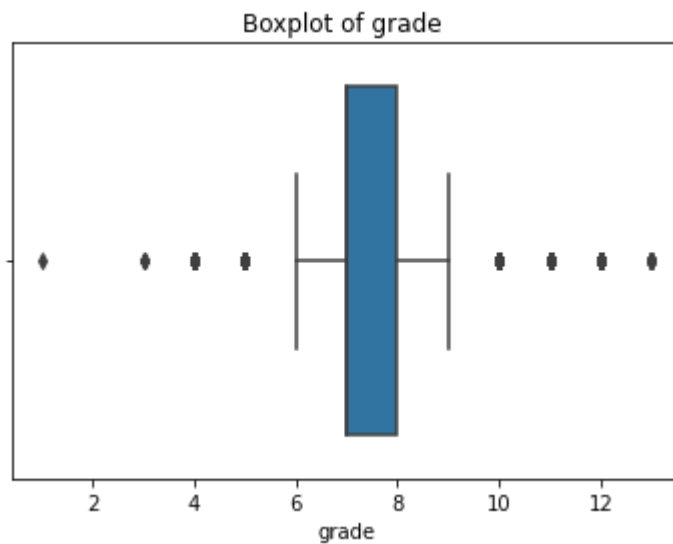
```
In [15]: # 3. box plot of floors to detect outliers
plt.title('Boxplot of floors ')
sns.boxplot(data=df,x='floors')
plt.show()
```



```
In [16]: # 4. box plot of sqft_living to detect outliers
plt.title('Boxplot of sqft_living')
sns.boxplot(data=df,x='sqft_living')
plt.show()
```



```
In [17]: # 5. box plot of grade to detect outliers
plt.title('Boxplot of grade ')
sns.boxplot(data=df,x='grade')
plt.show()
```



**More detail information of the boxplot:**

```
In [18]: df['grade'].describe()
```

```
Out[18]: count      21613.000000  
mean          7.656873  
std           1.175459  
min           1.000000  
25%           7.000000  
50%           7.000000  
75%           8.000000  
max           13.000000  
Name: grade, dtype: float64
```

```
In [19]: df['sqft_living'].describe()
```

```
Out[19]: count      21613.000000  
mean      2079.899736  
std       918.440897  
min       290.000000  
25%      1427.000000  
50%      1910.000000  
75%      2550.000000  
max     13540.000000  
Name: sqft_living, dtype: float64
```

```
In [20]: df['floors'].describe()
```

```
Out[20]: count      21613.000000  
mean          1.494309  
std           0.539989  
min           1.000000  
25%           1.000000  
50%           1.500000  
75%           2.000000  
max           3.500000  
Name: floors, dtype: float64
```

```
In [21]: df['bathrooms'].describe()
```

```
Out[21]: count      21613.000000  
mean          2.114757  
std           0.770163  
min           0.000000  
25%           1.750000  
50%           2.250000  
75%           2.500000  
max           8.000000  
Name: bathrooms, dtype: float64
```

```
In [22]: df['price_new'].describe()
```

```
Out[22]: count      21613.000000
mean         540.088142
std          367.127196
min           75.000000
25%          321.950000
50%          450.000000
75%          645.000000
max          7700.000000
Name: price_new, dtype: float64
```

***we are checking if there is any repetition on our data. We found out that the repetition mostly are in the variable of 'waterfront' and 'yr\_renovated'.***

However, we are not going to use those two variables.

```
In [23]: num_rows = len(df.index)
low_information_cols = []

for col in df.columns:
    cnts = df[col].value_counts(dropna=False)
    top_pct = (cnts/num_rows).iloc[0]

    if top_pct > 0.95:
        low_information_cols.append(col)
        print('{0}: {1:.5f}%'.format(col, top_pct*100))
        print(cnts)
        print()
```

```
waterfront: 99.24582%
0      21450
1       163
Name: waterfront, dtype: int64
```

```
yr_renovated: 95.77106%
0      20699
2014      91
2013      37
2003      36
2000      35
...
1934       1
1959       1
1951       1
1948       1
1944       1
Name: yr_renovated, Length: 70, dtype: int64
```

```
In [24]: df_dedupped = df.drop('id', axis=1).drop_duplicates()
print(df.shape)
print(df_dedupped.shape)
```

(21613, 18)

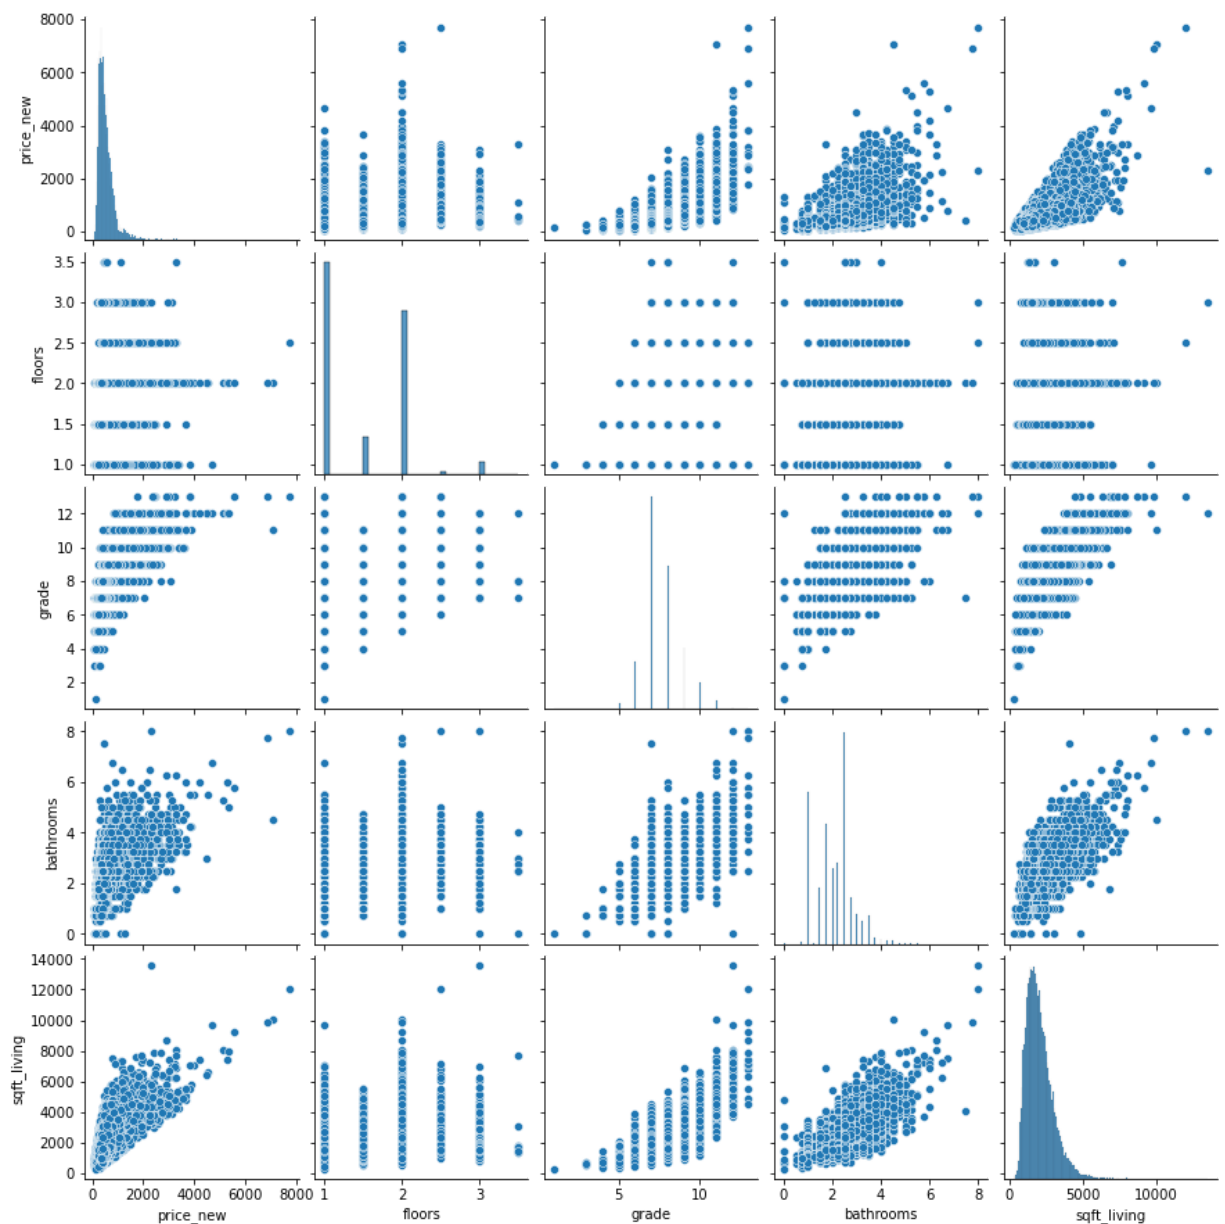
(21613, 17)

There is no duplicate house records in this dataset

## Pairplot

```
In [25]: sns.pairplot(df[['price_new', 'floors', 'grade', 'bathrooms', 'sqft_living']])
```

Out[25]: <seaborn.axisgrid.PairGrid at 0x119190e80>



## Rearession model

```
In [26]: import patsy
import statsmodels.api as sm
import statsmodels.formula.api as smf

#our dataframe is called df
df.head(5)
```

```
Out[26]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	water
0	7129300520	20141013T000000	221900	3	1.00	1180	5650	1.0	
1	6414100192	20141209T000000	538000	3	2.25	2570	7242	2.0	
2	5631500400	20150225T000000	180000	2	1.00	770	10000	1.0	
3	2487200875	20141209T000000	604000	4	3.00	1960	5000	1.0	
4	1954400510	20150218T000000	510000	3	2.00	1680	8080	1.0	

```
In [27]: #1. Specify the regression
#price=β0+β1grade+β2sqft_living+β3floors+β4bathrooms+ε. #this is a mathematical equation
#2. Create the model
#Using the statsmodel syntax, we have
#price ~ grade + sqft_living + floors + bathrooms

#WE are creating four different regression models
# #model 1
price_model1= smf.ols('price_new ~ grade', data=df) #running our main explanatory variable
# #model 2
price_model2= smf.ols('price_new ~ grade + sqft_living', data=df) #adding sqft_living
# #model 3
price_model3 = smf.ols('price_new ~ grade + sqft_living + floors', data=df) #adding floors
# #model 4
price_model4 = smf.ols('price_new ~ grade + sqft_living + floors + bathrooms', data=df)
```

```
In [28]: #use .fit() method to estimate the model fit for OLS
results1 = price_model1.fit() #model 1 fitting
results2 = price_model2.fit() #model 2 fitting
results3 = price_model3.fit() #model 3 fitting
results4 = price_model4.fit() #model 4 fitting
```

We print the summary and check the P values for all the independent variables to see if they are statistical significance.

```
In [29]: print(results1.summary())
print(results2.summary())
print(results3.summary())
print(results4.summary())
```

```

                                OLS Regression Results
=====
=
Dep. Variable:                  price_new    R-squared:                  0.44
5
Model:                          OLS        Adj. R-squared:          0.44
5
Method:                        Least Squares    F-statistic:                1.736e+0
4
Date:                          Sun, 04 Apr 2021    Prob (F-statistic):          0.0
0
Time:                          11:35:34        Log-Likelihood:             -1.5194e+0
5
No. Observations:              21613          AIC:                       3.039e+0
5
Df Residuals:                  21611          BIC:                       3.039e+0
5
Df Model:                      1
Covariance Type:               nonrobust

```

We reorganized the summary from above and put important results for all models into one summary table.

```
In [30]: from statsmodels.iolib.summary2 import summary_col #create regression table

table = summary_col(
    [results1, results2, results3, results4],
    model_names = ['Model 1', 'Model 2', 'Model 3', 'Model 4'],
    stars=True, #level of significance
    regressor_order = ['Intercept', 'grade', 'sqft_living', 'floors', 'bathroom'],
    float_format='%0.2f',
    drop_omitted = False #
)

table.add_title('Does higher grade have a higher housing price?') # RQ: How does
# controlling sqft_living, floors, and bathrooms

print(table.as_text())

# The 'w' is for 'write'.
fout = open('table1.txt', 'w')
fout.write(table.as_text())
fout.close()
```

```
Does higher grade have a higher housing price?
=====
```

	Model 1	Model 2	Model 3	Model 4
Intercept	-1056.04*** (12.26)	-598.11*** (13.30)	-602.58*** (13.26)	-601.35*** (13.24)
grade	208.46*** (1.58)	98.55*** (2.24)	107.64*** (2.35)	109.45*** (2.36)
sqft_living		0.18*** (0.00)	0.18*** (0.00)	0.20*** (0.00)
floors			-43.96*** (3.54)	-34.65*** (3.76)
bathrooms				-26.67*** (3.65)
R-squared	0.45	0.53	0.54	0.54
R-squared Adj.	0.45	0.53	0.54	0.54

```
=====
Standard errors in parentheses.
* p<.1, ** p<.05, ***p<.01
```